

Data Analytics Lab: Assignment - 3

A Mathematical Essay on Naive Bayes Classifier

Nagappan N

Department of Metallurgical and Materials Engineering

IIT Madras

Chennai, India

mm19b040@smail.iitm.ac.in

Abstract—The aim of this project is to predict whether or not a person earns more than \$50k per year. We have used the Bernoulli Naive Bayes classifier to achieve this. Our model is able to predict with an accuracy of 79.85%.

I. INTRODUCTION

This project mainly aims to predict whether or not a person earns more than \$50k per year. We try to see what are the factors that results in more earnings for a person. We use the 1994 Census bureau database by Ronny Kohavi and Barry Becker for this purpose.

The "naive" assumption that each pair of features is conditionally independent given the value of the class variable underlies a class of supervised learning algorithms collectively referred to as "naive Bayes methods." Bernoulli Naive Bayes is used for data that is distributed according to multivariate Bernoulli distributions.

Using the naive bayes classifier, we predict whether or not a person earns more than \$50k per year. We use the input features of age, person's current job related information, educational status, marital status, sex, race and native country for this prediction. The model developed would help us in making some crucial inferences. A similar version of this model can be used a way to classify below above and below the poverty line. Hence, this kind of model would be very handy for government based agencies and NGOs.

This paper starts with the description of naive bayes classifier. It is followed by the description of the datasets. This includes data visualization and processing. The following section includes details about how the model has been implemented. Finally, we conclude with the key inferences from our project.

II. NAIVE BAYES CLASSIFIER

A group of supervised learning algorithms known as naive Bayes methods utilise Bayes' theorem with the "naive" assumption that each pair of features is conditionally independent given the value of the class variable. Given the class variable y and the dependent feature vectors x_1 through x_n , the Bayes theorem describes the relationship in

the following way.

$$P(y | x_1, \dots, x_n) = \frac{P(y)P(x_1, \dots, x_n | y)}{P(x_1, \dots, x_n)}$$

$$\hat{y} = \arg \max_y P(y) \prod_{i=1}^n P(x_i | y),$$

Then we may estimate $P(y)$ and $P(x_i|y)$ using Maximum A Posteriori (MAP) estimation; the former is then the relative frequency of class y in the training set.

The assumptions that different naive Bayes classifiers make about the distribution of $P(x_i|y)$ are what distinguish them from one another:

- Gaussian Naive Bayes: GaussianNB implements the Gaussian Naive Bayes algorithm for classification. The likelihood of the features is assumed to be Gaussian.
- Multinomial Naive Bayes: MultinomialNB implements the naive Bayes algorithm for multinomially distributed data, The parameters y is estimated by a smoothed version of maximum likelihood.
- Complement Naive Bayes: CNB is an adaptation of the standard multinomial naive Bayes algorithm for imbalanced data sets. CNB uses statistics from the complement of each class to compute the model's weights.
- Categorical Naive Bayes: CategoricalNB implements the categorical naive Bayes algorithm for categorically distributed data. It assumes that each feature, which is described by the index i , has its own categorical distribution.
- Bernoulli Naive Bayes: BernoulliNB is used for data that is distributed according to multivariate Bernoulli distributions; i.e., there may be multiple features but each one is assumed to be a binary-valued (Bernoulli, boolean) variable. The expression for this is as follows:

$$P(x_i | y) = P(x_i = 1 | y)x_i + (1 - P(x_i = 1 | y))(1 - x_i)$$

III. DATASETS

The dataset has the following columns:

- 1) age: The age of the person. It is a continuous variable.
- 2) work class: This is a categorical variable describing the working class of the person. It takes 8 values: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.
- 3) fnlwgt:
- 4) education: The level of education obtained by the person. This is again a categorical variable which takes one of the 16 values: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
- 5) education-num: Number of years of education pursued by the person.
- 6) marital status: The marital status of the person is also a categorical variable which takes one of the following values: marital-status Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
- 7) Occupation: This is also a categorical variable describing the occupation of the person. It takes one of these values: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces
- 8) relationship: This is defined as a categorical variable that can take the values: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
- 9) race: This categorical variables describes the race of the person which can be White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other or Black.
- 10) sex: The sex of the person.
- 11) capital gain
- 12) capital loss
- 13) Hours-per-week: The number of hours a person works in a week is given by this variable.
- 14) Native-country: The country that the person belongs to.
- 15) salary: Tells us if the person earns greater than \$50k or lesser than \$50k per year.




Fig. 1. Heatmap showing the distribution of missing values in our dataset

The employment variables work-class, occupation, and native-country all have missing values. We decide to add a distinct category for missing data as "Unknown" because the percentage of missing values is larger.

We can now clean the data and discretize the continuous variables of age, hours-per-week, education count, capital loss, and capital gain into categorical ones based on the relative density distributions of their salary class since we have the highest number of input variables of the categorical type. In order to condense the enormous number of classes, we also examine each categorical variable's relative density distribution for its respective salary class.

- Work class: Here, we consolidate all federal, local, and state government-related classes. We also eliminate redundancy by integrating the never-worked and unpaid job classes. Additionally, as the pay scales for the self-employed and the private sector are comparable, we have combined the two.




Fig. 2. Percentage of people who earn above and below \$50k per year based on their work class.

A. Data Processing

First, we visualize the missing values in the dataset.

- Number of years of education: We can see a clear difference after 13 years of education in the salary fractions this is a direct correlation between persons who completed schooling and people who did not.




Fig. 3. Percentage of people who earn above and below \$50k per year based on number of years of education obtained by the person.




Fig. 6. Percentage of people who earn above and below \$50k per year based on the person's sex.

- Education: In this column we can combine whole of schooling together and also combine Bachelor and master degree holders. We can also see that doctoral and professional school education levels are comparable, thus we can combine the two. HS graduates and other college groups are also able to be combined. Additionally, an Assoc category may combine the Assoc-voc and Assoc-acdm categories.




Fig. 4. Percentage of people who earn above and below \$50k per year based on their education.

- Age: Based on the following plot we discretize age into three categories in intervals of 40 and as the fraction of salaries are similar in the first and third categories we combine them.



Fig. 7. Percentage of people who earn above and below \$50k per year based on the person's age.

- Race: We see white and Asians have higher salary groups than other races.




Fig. 5. Percentage of people who earn above and below \$50k per year based on the person's race.

- Sex: We see that males have higher salary fraction than females




Fig. 8. Histogram plot using polynomial element for the variable age.

- Hours per week count plot: We can see from the probability graph that we can discretize the domain into 0-50 and 50-100.




Fig. 9. Histogram plot using polynomial element for the variable hours per week.




Fig. 12. Percentage of people who earn above and below \$50k per year based on capital gain.




Fig. 13. Histogram plot using step element for the variable capital loss.

Fig. 10. Percentage of people who earn above and below \$50k per year based on number of hours the person works.

- Capital gain and capital loss: Using probability density charts, we discretize and combine the comparable categories in a manner similar to the aforementioned continuous plots.




Fig. 14. Percentage of people who earn above and below \$50k per year based on capital loss.

- Marital status: We can cut down different number of categories into just two With / without spouse.




Fig. 11. Histogram plot using step element for the variable capital gain.





Fig. 15. Percentage of people who earn above and below \$50k per year based on marital status.



- **Occupation:** Based on the wage distributions, we can divide many professions into two main groups: high and low income groups.




Fig. 16. Percentage of people who earn above and below \$50k per year based on occupation.

- **Relationship:** Here again each category can be brought into whether or not they are in a family.




Fig. 17. Percentage of people who earn above and below \$50k per year based on relationship.

- **Native Country:** We can categorise into US and non-US countries because the majority of data is in the US.




Fig. 18. Percentage of people who earn above and below \$50k per year based on native country.

IV. MODELLING

We only use columns that, after cleaning, may be divided into only two groups. We drop the work class and education columns as they essentially cover the same information as the occupation and education number columns. So, now, a

Bernoulli Naive Bayes classifier model is appropriate.

Before using BernoulliNB to train and forecast the data, we must first use Label Encoder to assign a 0 or 1 to each category. To assess performance following training, we separated the data into train and test datasets.

We may compare the target salary variable's predicted and actual values by utilising the accuracy score as the performance metric. Our model is able to make predictions with an accuracy of 79.85%

V. CONCLUSIONS

We find that a simple model like naive bayes is good to make predictions for real world problems like this. We have implemented a Bernoulli Naive Bayes classifier on the dataset to predict whether or not a person earns more than \$50k per year with an accuracy of 79.85%.

EE4708: Data Analytics Lab

Assignment 3

Name: Nagappan N

Roll number: MM19B040

Importing libraries

In [75]:

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
```

Understanding the Dataset

In [76]:

```
df=pd.read_excel('adult.xlsx')
```

In [77]:

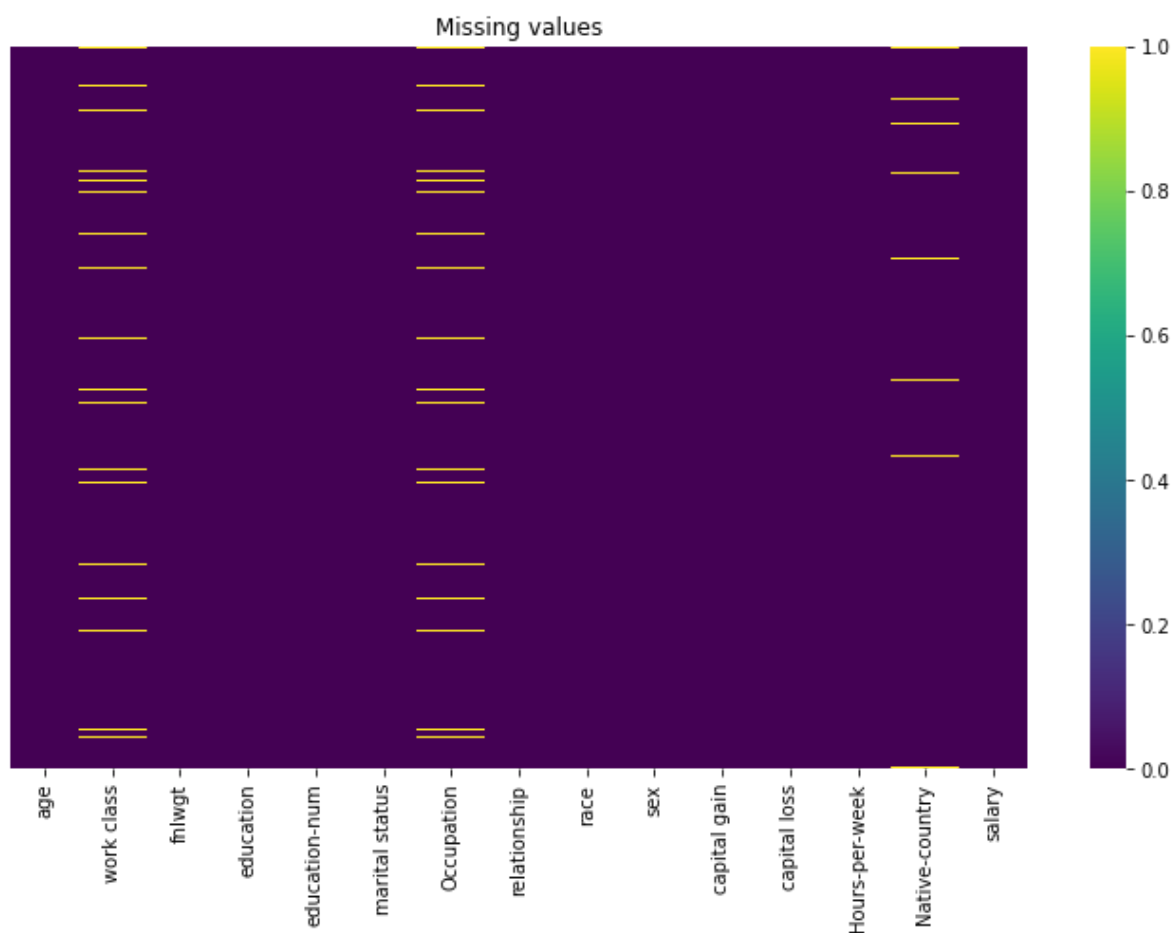
```
df.head()
```

Out[77]:

	age	work class	fnlwgt	education	education-num	marital status	Occupation	relationship	race	sex
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female

In [78]:

```
df.replace('\?', np.nan, regex=True, inplace=True)
plt.figure(figsize=[12,7])
sns.heatmap(df.isnull(), cmap='viridis')
plt.yticks(ticks=[])
plt.title('Missing values')
plt.savefig('Missingvalues.png')
```



As missing values are not easily determinable we create a new category unknown for all missing values alike

In [79]:

```
df.fillna('Unknown',inplace=True)
```

In [80]:

```
df['work class']=df['work class'].astype('category')
df['education']=df['education'].astype('category')
df['marital status']=df['marital status'].astype('category')
df['Occupation']=df['Occupation'].astype('category')
df['relationship']=df['relationship'].astype('category')
df['race']=df['race'].astype('category')
df['sex']=df['sex'].astype('category')
df['Native-country']=df['Native-country'].astype('category')
df['salary']=df['salary'].astype('category')
```

In [81]:



```
df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    32561 non-null  int64
1   work class             32561 non-null  category
2   fnlwgt                 32561 non-null  int64
3   education              32561 non-null  category
4   education-num          32561 non-null  int64
5   marital status         32561 non-null  category
6   Occupation             32561 non-null  category
7   relationship           32561 non-null  category
8   race                   32561 non-null  category
9   sex                    32561 non-null  category
10  capital gain           32561 non-null  int64
11  capital loss           32561 non-null  int64
12  Hours-per-week         32561 non-null  int64
13  Native-country         32561 non-null  category
14  salary                 32561 non-null  category
dtypes: category(9), int64(6)
memory usage: 1.8 MB
```

We next look into how the distribution of few of the continuous features in our dataset.

In [82]:




```
plt.figure(figsize=(10,7))
sns.histplot(df['age'])
font2 = {'family':'serif','color':'black','size':12}
font1 = {'family':'serif','color':'black','size':14}
plt.title('Age Distribution',fontdict=font1)
plt.ylabel('Count',fontdict=font2)
plt.xlabel('Age',fontdict=font2)
plt.savefig('agedist.png',dpi=2048)
```



In [83]:



```
plt.figure(figsize=(10,7))
sns.histplot(df['Hours-per-week'])
font2 = {'family':'serif','color':'black','size':12}
font1 = {'family':'serif','color':'black','size':14}
plt.title('Age Distribution',fontdict=font1)
plt.ylabel('Count',fontdict=font2)
plt.xlabel('Hours per week',fontdict=font2)
plt.savefig('hoursdist.png',dpi=2048)
```



As we have maximum categorical variables we need to bring all data in categorical form

Now we proceed to clean and extract the features from the data by following steps

- If it is a continuous variable we discretize and then combine discretized categories of similar income groups to have not more than 2 categories
- If it is a categorical variable we clean it to the best possible extent and use the ones which has only two categories into a bernoulli Naive bayes algorithm.

In [84]:



```
df.groupby('work class')['salary'].value_counts(normalize=True)
```

Out[84]:

work class	salary	
Federal-gov	<=50K	0.613542
	>50K	0.386458
Local-gov	<=50K	0.705208
	>50K	0.294792
Never-worked	<=50K	1.000000
	>50K	0.000000
Private	<=50K	0.781327
	>50K	0.218673
Self-emp-inc	>50K	0.557348
	<=50K	0.442652
Self-emp-not-inc	<=50K	0.715073
	>50K	0.284927
State-gov	<=50K	0.728043
	>50K	0.271957
Without-pay	<=50K	1.000000
	>50K	0.000000
Unknown	<=50K	0.895969
	>50K	0.104031

Name: salary, dtype: float64

Lets map all government jobs together and combine private and self-emp-not-inc with private, and also never worked and without pay together to create better categories

In [85]:

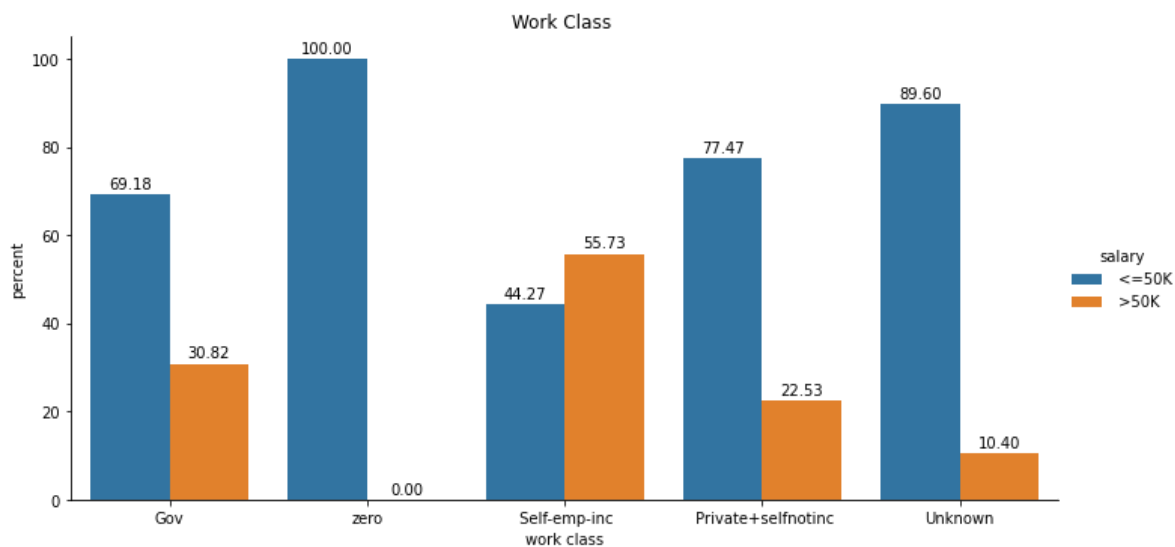


```
mappings={' Federal-gov':'Gov', ' Local-gov':'Gov', ' State-gov':'Gov', ' Never-worked':'zero'}
df.replace(mappings,inplace=True)
```

In [86]:

```
x='work class'
y='salary'
dfp=(df
.groupby(x)[y]
.value_counts(normalize=True)
.mul(100)
.rename('percent')
.reset_index())
plt.figure(figsize=[15,7])
g=sns.catplot(x='work class',y='percent',hue=y,data=dfp,kind='bar',height=5,aspect=2)
for container in g.ax.containers:
    g.ax.bar_label(container, fmt='%.2f', padding=2)
plt.title('Work Class ')
plt.savefig('workclasscplot2.png')
```

<Figure size 1080x504 with 0 Axes>



In [87]:



```
df.groupby('education-num')['salary'].value_counts(normalize=True)
```

Out[87]:

education-num	salary	
1	<=50K	1.000000
	>50K	0.000000
2	<=50K	0.964286
	>50K	0.035714
3	<=50K	0.951952
	>50K	0.048048
4	<=50K	0.938080
	>50K	0.061920
5	<=50K	0.947471
	>50K	0.052529
6	<=50K	0.933548
	>50K	0.066452
7	<=50K	0.948936
	>50K	0.051064
8	<=50K	0.923788
	>50K	0.076212
9	<=50K	0.840491
	>50K	0.159509
10	<=50K	0.809765
	>50K	0.190235
11	<=50K	0.738784
	>50K	0.261216
12	<=50K	0.751640
	>50K	0.248360
13	<=50K	0.585247
	>50K	0.414753
14	>50K	0.556587
	<=50K	0.443413
15	>50K	0.734375
	<=50K	0.265625
16	>50K	0.740920
	<=50K	0.259080

Name: salary, dtype: float64

In [88]:




```
df['education-num']=pd.cut(df['education-num'],bins=[0,12,16],labels=['0-8','13-16'])
```

In [89]:



```
x1='education-num'
y='salary'
dfp=(df
.groupby(x1)[y]
.value_counts(normalize=True)
.mul(100)
.rename('percent')
.reset_index())
plt.figure(figsize=[15,7])
g=sns.catplot(x=x1,y='percent',hue=y,data=dfp,kind='bar',height=5,aspect=2)
for container in g.ax.containers:
    g.ax.bar_label(container, fmt='%.2f', padding=2)
plt.title('Number of years of education')
plt.savefig('edunumcplot.png')
```

<Figure size 1080x504 with 0 Axes>



In [90]:

```
df.groupby('education')['salary'].value_counts(normalize=True)
```

Out[90]:

education	salary	
10th	<=50K	0.933548
	>50K	0.066452
11th	<=50K	0.948936
	>50K	0.051064
12th	<=50K	0.923788
	>50K	0.076212
1st-4th	<=50K	0.964286
	>50K	0.035714
5th-6th	<=50K	0.951952
	>50K	0.048048
7th-8th	<=50K	0.938080
	>50K	0.061920
9th	<=50K	0.947471
	>50K	0.052529
Assoc-acdm	<=50K	0.751640
	>50K	0.248360
Assoc-voc	<=50K	0.738784
	>50K	0.261216
Bachelors	<=50K	0.585247
	>50K	0.414753
Doctorate	>50K	0.740920
	<=50K	0.259080
HS-grad	<=50K	0.840491
	>50K	0.159509
Masters	>50K	0.556587
	<=50K	0.443413
Preschool	<=50K	1.000000
	>50K	0.000000
Prof-school	>50K	0.734375
	<=50K	0.265625
Some-college	<=50K	0.809765
	>50K	0.190235

Name: salary, dtype: float64


In [91]:

```
mappings={' 10th':' School',' 11th':' School',' 12th':' School',' 1st-4th':' School',' 5th-  
df.replace(mappings,inplace=True)
```

In [92]:

```
x1='education'
y='salary'
dfp=(df
.groupby(x1)[y]
.value_counts(normalize=True)
.mul(100)
.rename('percent')
.reset_index())
plt.figure(figsize=[15,7])
g=sns.catplot(x=x1,y='percent',hue=y,data=dfp,kind='bar',height=5,aspect=2)
for container in g.ax.containers:
    g.ax.bar_label(container, fmt='%.2f', padding=2)
plt.savefig('educplot.png')
```

<Figure size 1080x504 with 0 Axes>



In [93]:



```
df.groupby('race')['salary'].value_counts(normalize=True)
```

Out[93]:

race	salary	
Amer-Indian-Eskimo	<=50K	0.884244
	>50K	0.115756
Asian-Pac-Islander	<=50K	0.734360
	>50K	0.265640
Black	<=50K	0.876120
	>50K	0.123880
Other	<=50K	0.907749
	>50K	0.092251
White	<=50K	0.744140
	>50K	0.255860

Name: salary, dtype: float64

As we can see the fraction of greater salary holders is similar among white and asians and all other categories have also similar lower fraction of >50K salary.

In [94]:



```
mappings={' Amer-Indian-Eskimo':' Other',' Asian-Pac-Islander':'White+Asian',' Black':' Oth
df['race'].replace(mappings,inplace=True)
df.groupby('race')['salary'].value_counts(normalize=True)
```

Out[94]:

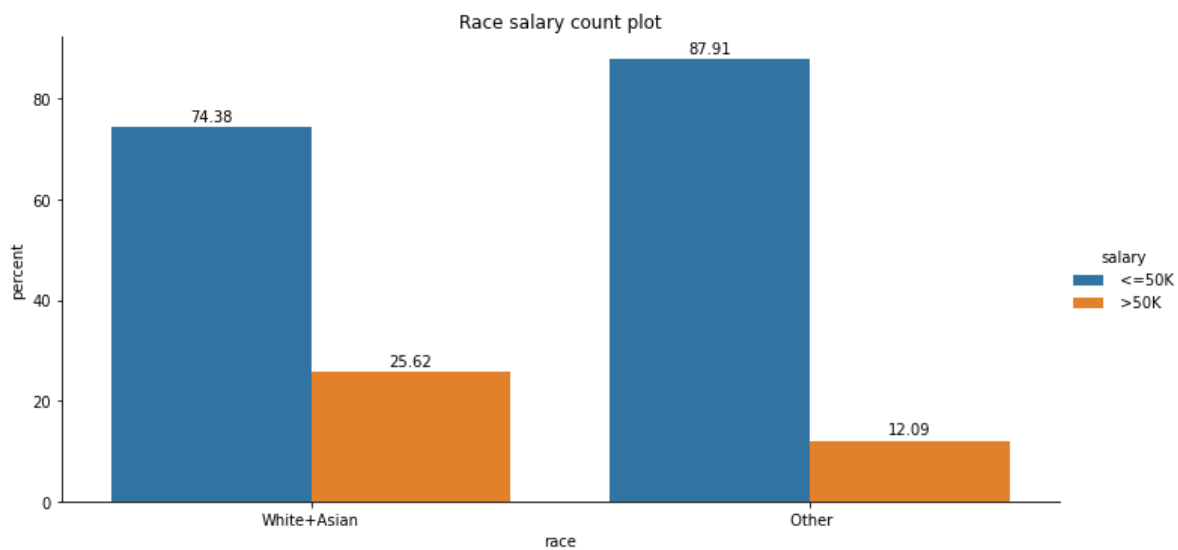
race	salary	
White+Asian	<=50K	0.743788
	>50K	0.256212
Other	<=50K	0.879115
	>50K	0.120885

Name: salary, dtype: float64

In [95]:

```
x1='race'
y='salary'
dfp=(df
.groupby(x1)[y]
.value_counts(normalize=True)
.mul(100)
.rename('percent')
.reset_index())
plt.figure(figsize=[15,7])
g=sns.catplot(x=x1,y='percent',hue=y,data=dfp,kind='bar',height=5,aspect=2)
for container in g.ax.containers:
    g.ax.bar_label(container, fmt='%.2f', padding=2)
plt.title('Race salary count plot')
plt.savefig('racecplot2.png')
```

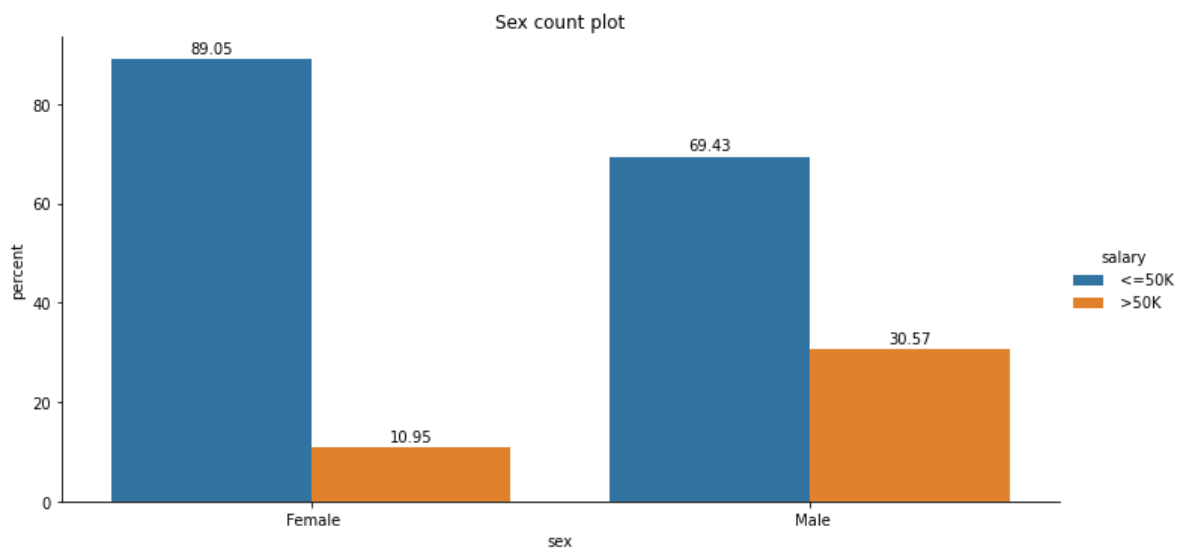
<Figure size 1080x504 with 0 Axes>



In [96]:

```
x1='sex'
y='salary'
dfp=(df
.groupby(x1)[y]
.value_counts(normalize=True)
.mul(100)
.rename('percent')
.reset_index())
plt.figure(figsize=[15,7])
g=sns.catplot(x=x1,y='percent',hue=y,data=dfp,kind='bar',height=5,aspect=2)
for container in g.ax.containers:
    g.ax.bar_label(container, fmt='%.2f', padding=2)
plt.title('Sex count plot')
plt.savefig('sexcplot.png')
```

<Figure size 1080x504 with 0 Axes>



In [97]:

```
plt.figure(figsize=[10,5])
sns.histplot(data=df,x='age',hue='salary',multiple='dodge',element='poly',fill=False,stat='probability')
plt.title('Age salary count plot')
plt.savefig('Agecplot1.png')
```



Similarly we can discretize the age and hours per week variable as intervals according to the above probability distribution

In [98]:

```
df['age']=pd.cut(df.age,bins=[i for i in range(0,160,40)],labels=[str(i)+'-'+str(i+40) for
```


In [99]:

```
mapping={'80-120':'0-40 & 80-90','0-40':'0-40 & 80-90'}
df['age'].replace(mapping,inplace=True)
```

In [100]:

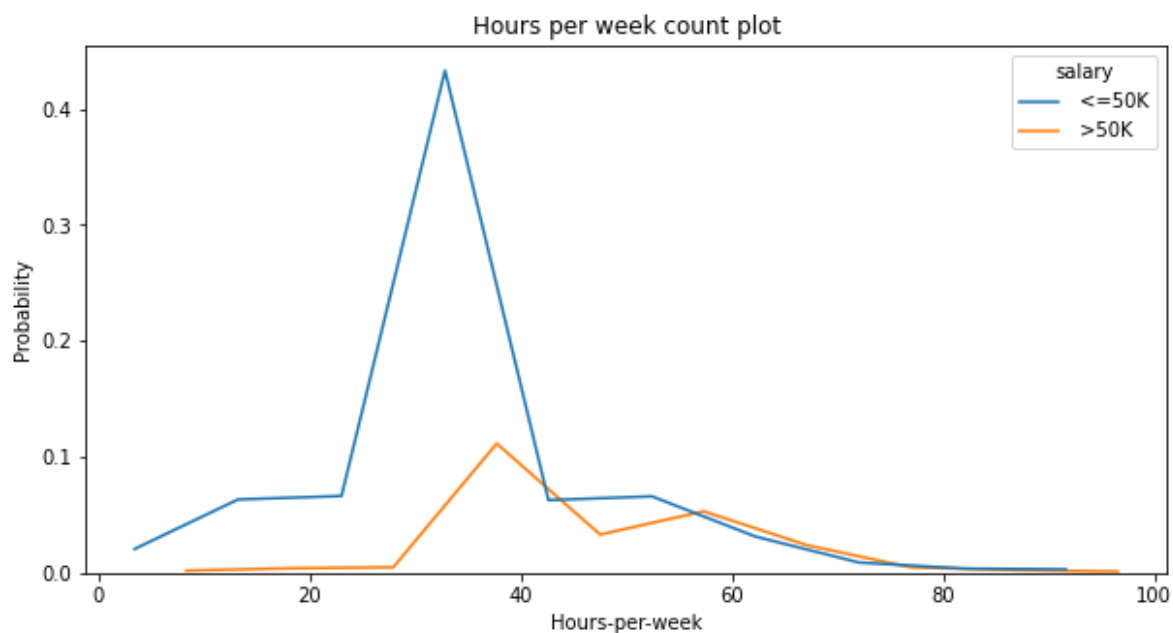
```
x1='age'
y='salary'
dfp=(df
.groupby(x1)[y]
.value_counts(normalize=True)
.mul(100)
.rename('percent')
.reset_index())
plt.figure(figsize=[15,7])
g=sns.catplot(x=x1,y='percent',hue=y,data=dfp,kind='bar',height=5,aspect=2)
for container in g.ax.containers:
    g.ax.bar_label(container, fmt='%.2f', padding=2)
plt.title('Age count plot')
plt.savefig('agecplot2.png')
```

<Figure size 1080x504 with 0 Axes>



In [101]:

```
plt.figure(figsize=[10,5])
sns.histplot(data=df,x='Hours-per-week',hue='salary',bins=10,multiple='dodge',element='poly')
plt.title('Hours per week count plot')
plt.savefig('hourscplot.png')
```



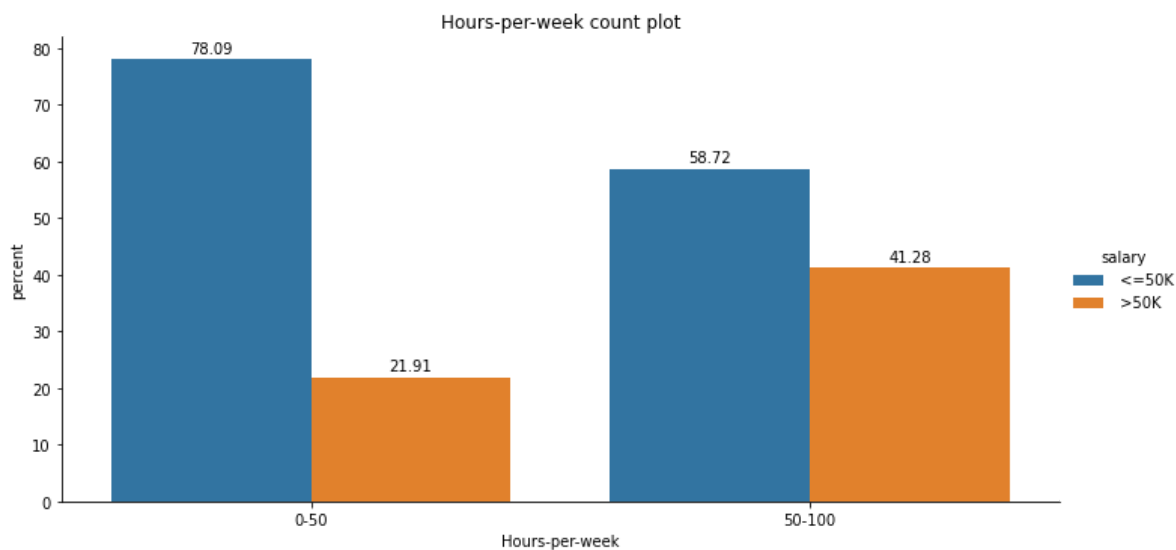
In [102]:

```
df['Hours-per-week']=pd.cut(df['Hours-per-week'],bins=[0,50,100],labels=['0-50','50-100'])
```

In [103]:

```
x1='Hours-per-week'
y='salary'
dfp=(df
.groupby(x1)[y]
.value_counts(normalize=True)
.mul(100)
.rename('percent')
.reset_index())
plt.figure(figsize=[15,7])
g=sns.catplot(x=x1,y='percent',hue=y,data=dfp,kind='bar',height=5,aspect=2)
for container in g.ax.containers:
    g.ax.bar_label(container, fmt='%.2f', padding=2)
plt.title('Hours-per-week count plot')
plt.savefig('Hourscplot2.png')
```

<Figure size 1080x504 with 0 Axes>



Now let us look at the capital gain and capital loss distributions and proceed to discretize the same

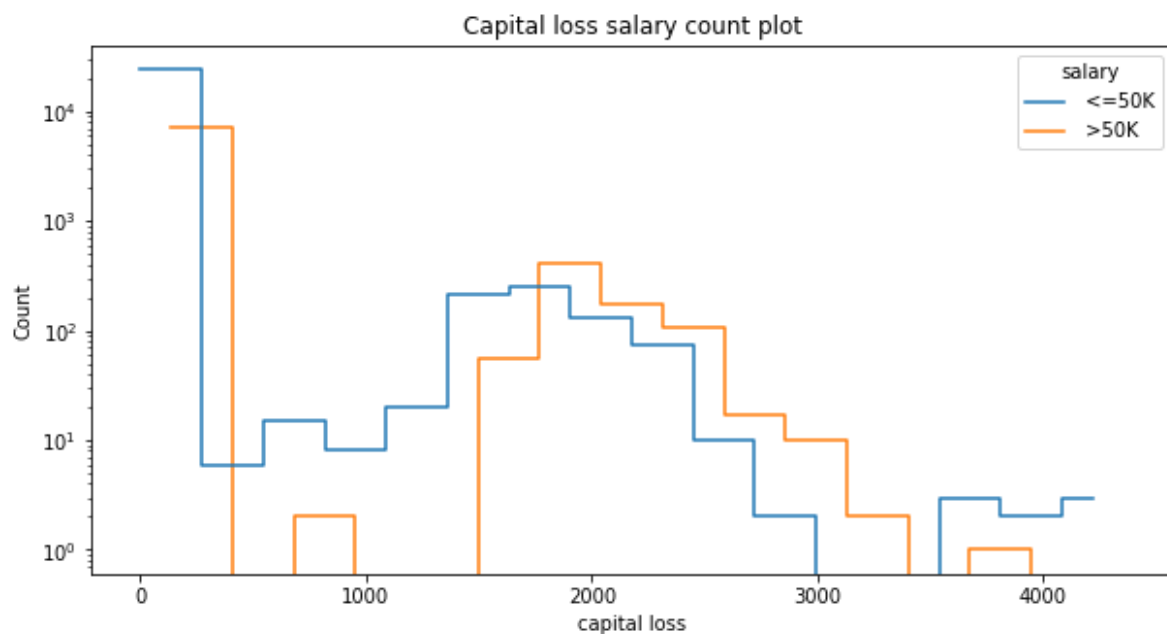
In [104]:

```
plt.figure(figsize=[10,5])
sns.histplot(data=df,x='capital gain',hue='salary',multiple='dodge',element='step',fill=False)
plt.title('Capital gain salary count plot')
plt.savefig('Capitalgain.png')
```



In [105]:

```
plt.figure(figsize=[10,5])
sns.histplot(data=df,x='capital loss',hue='salary',multiple='dodge',element='step',fill=False)
plt.title('Capital loss salary count plot')
plt.savefig('Capitalloss.png')
```



In [106]:

```
cap_loss_bins=[0,1000,4000,5000]
cap_gain_bins=[0,5000,100000]
cap_loss_labels=[str(cap_loss_bins[i])+'-'+str(cap_loss_bins[i+1]) for i in range(len(cap_l
cap_gain_labels=[str(cap_gain_bins[i])+'-'+str(cap_gain_bins[i+1]) for i in range(len(cap_g
```

In [107]:

```
df['capital gain']=pd.cut(df['capital gain'],bins=cap_gain_bins,labels=cap_gain_labels,incl
```

In [108]:

```
df['capital loss']=pd.cut(df['capital loss'],bins=cap_loss_bins,labels=cap_loss_labels,incl
```


In [109]:

```
mappings={'4000-5000':'0-1000 & 4000-5000','0-1000':'0-1000 & 4000-5000'}
df['capital loss'].replace(mappings,inplace=True)
```

In [110]:

```
x1='capital gain'
y='salary'
dfp=(df
.groupby(x1)[y]
.value_counts(normalize=True)
.mul(100)
.rename('percent')
.reset_index())
plt.figure(figsize=[15,7])
g=sns.catplot(x=x1,y='percent',hue=y,data=dfp,kind='bar',height=5,aspect=2)
for container in g.ax.containers:
    g.ax.bar_label(container, fmt='%.2f', padding=2)
plt.title('Capital gain salary count plot')
plt.savefig('Capitalgain2.png')
```

<Figure size 1080x504 with 0 Axes>

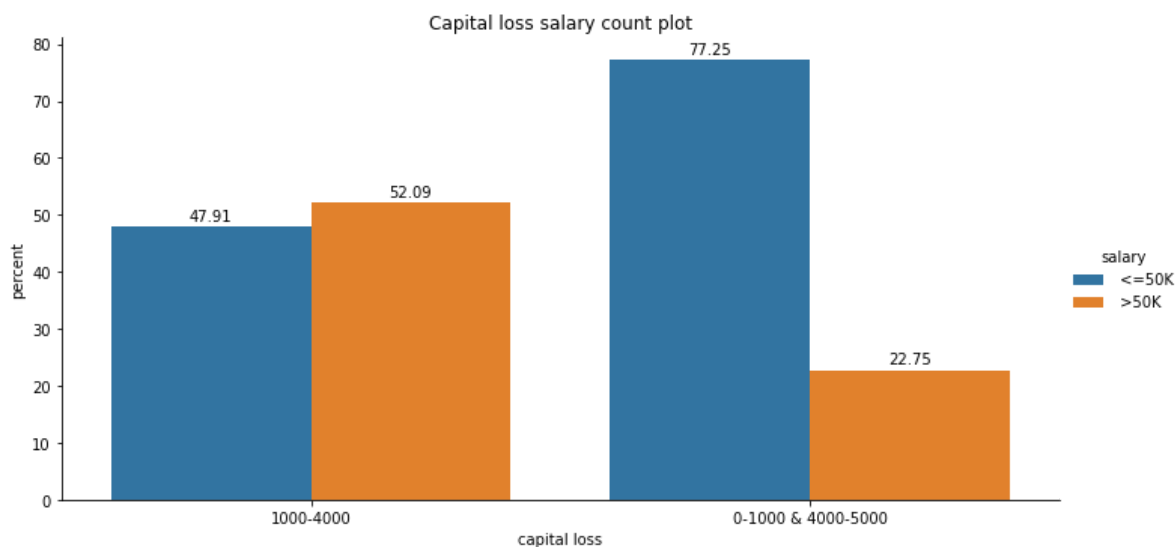


In [111]:



```
x1='capital loss'
y='salary'
dfp=(df
.groupby(x1)[y]
.value_counts(normalize=True)
.mul(100)
.rename('percent')
.reset_index())
plt.figure(figsize=[15,7])
g=sns.catplot(x=x1,y='percent',hue=y,data=dfp,kind='bar',height=5,aspect=2)
for container in g.ax.containers:
    g.ax.bar_label(container, fmt='%.2f', padding=2)
plt.title('Capital loss salary count plot')
plt.savefig('Capitalloss2.png')
```

<Figure size 1080x504 with 0 Axes>



In [112]:

```
df.groupby('marital status')['salary'].value_counts(normalize=True)
```

Out[112]:

marital status	salary	
Divorced	<=50K	0.895791
	>50K	0.104209
Married-AF-spouse	<=50K	0.565217
	>50K	0.434783
Married-civ-spouse	<=50K	0.553152
	>50K	0.446848
Married-spouse-absent	<=50K	0.918660
	>50K	0.081340
Never-married	<=50K	0.954039
	>50K	0.045961
Separated	<=50K	0.935610
	>50K	0.064390
Widowed	<=50K	0.914401
	>50K	0.085599

Name: salary, dtype: float64

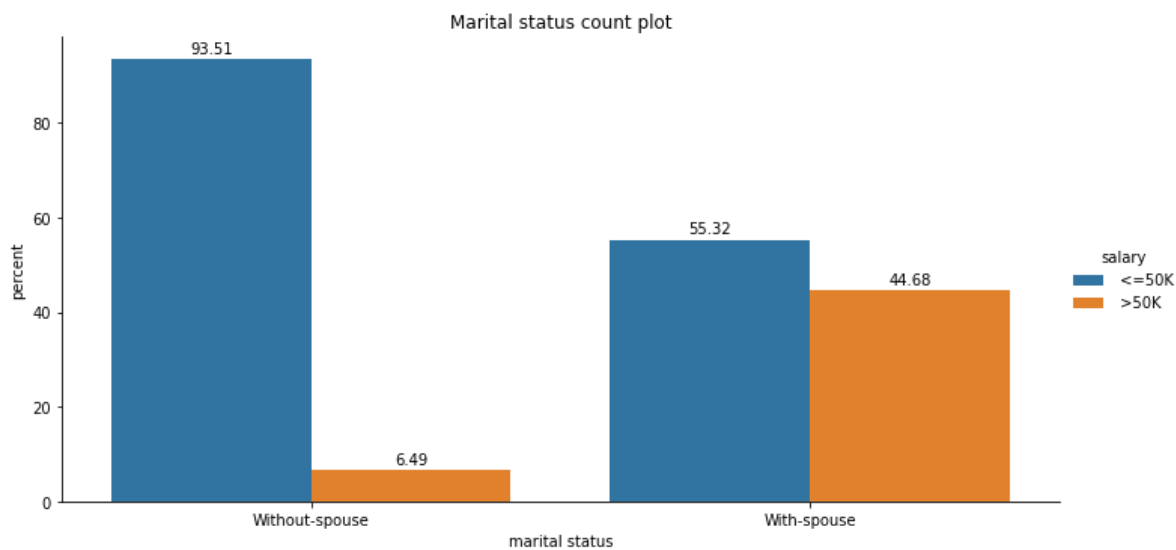
In [113]:

```
mapping={' Divorced': 'Without-spouse', ' Married-AF-spouse': 'With-spouse', ' Married-civ-spouse': 'With-spouse'}
df['marital status'].replace(mapping, inplace=True)
```

In [114]:

```
x1='marital status'
y='salary'
dfp=(df
.groupby(x1)[y]
.value_counts(normalize=True)
.mul(100)
.rename('percent')
.reset_index())
plt.figure(figsize=[15,7])
g=sns.catplot(x=x1,y='percent',hue=y,data=dfp,kind='bar',height=5,aspect=2)
for container in g.ax.containers:
    g.ax.bar_label(container, fmt='%.2f', padding=2)
plt.title('Marital status count plot')
plt.savefig('maritalcplot.png')
```

<Figure size 1080x504 with 0 Axes>



In [115]:

```
df.groupby('Occupation')['salary'].value_counts(normalize=True)
```

Out[115]:

Occupation	salary	
Adm-clerical	<=50K	0.865517
	>50K	0.134483
Armed-Forces	<=50K	0.888889
	>50K	0.111111
Craft-repair	<=50K	0.773359
	>50K	0.226641
Exec-managerial	<=50K	0.515986
	>50K	0.484014
Farming-fishing	<=50K	0.884306
	>50K	0.115694
Handlers-cleaners	<=50K	0.937226
	>50K	0.062774
Machine-op-inspct	<=50K	0.875125
	>50K	0.124875
Other-service	<=50K	0.958422
	>50K	0.041578
Priv-house-serv	<=50K	0.993289
	>50K	0.006711
Prof-specialty	<=50K	0.550966
	>50K	0.449034
Protective-serv	<=50K	0.674884
	>50K	0.325116
Sales	<=50K	0.730685
	>50K	0.269315
Tech-support	<=50K	0.695043
	>50K	0.304957
Transport-moving	<=50K	0.799624
	>50K	0.200376
Unknown	<=50K	0.896365
	>50K	0.103635

Name: salary, dtype: float64


In [116]:

```
mappings={' Adm-clerical':'Low-income-group', ' Armed-Forces':'Low-income-group', ' Farming-f  
df['Occupation'].replace(mappings,inplace=True)
```

In [117]:

```
x1='Occupation'
y='salary'
dfp=(df
.groupby(x1)[y]
.value_counts(normalize=True)
.mul(100)
.rename('percent')
.reset_index())
plt.figure(figsize=[15,7])
g=sns.catplot(x=x1,y='percent',hue=y,data=dfp,kind='bar',height=5,aspect=2)
for container in g.ax.containers:
    g.ax.bar_label(container, fmt='%.2f', padding=2)
plt.title('Occupation count plot')
plt.savefig('occcplot.png')
```

<Figure size 1080x504 with 0 Axes>



In [118]:



```
df.groupby('relationship')['salary'].value_counts(normalize=True)
```

Out[118]:

relationship	salary	
Husband	<=50K	0.551429
	>50K	0.448571
Not-in-family	<=50K	0.896930
	>50K	0.103070
Other-relative	<=50K	0.962283
	>50K	0.037717
Own-child	<=50K	0.986780
	>50K	0.013220
Unmarried	<=50K	0.936738
	>50K	0.063262
Wife	<=50K	0.524872
	>50K	0.475128

Name: salary, dtype: float64

In [119]:




```
mappings={' Husband':'Family',' Wife':'Family',' Other-relative':' Not-in-family',' Unmarri  
df['relationship'].replace(mappings,inplace=True)
```



In [120]:

```
x1='relationship'
y='salary'
dfp=(df
.groupby(x1)[y]
.value_counts(normalize=True)
.mul(100)
.rename('percent')
.reset_index())
plt.figure(figsize=[15,7])
g=sns.catplot(x=x1,y='percent',hue=y,data=dfp,kind='bar',height=5,aspect=2)
for container in g.ax.containers:
    g.ax.bar_label(container, fmt='%.2f', padding=2)
plt.title('relationships-count plot')
plt.savefig('relcplot.png')
```

<Figure size 1080x504 with 0 Axes>



In [121]:



```
df['Native-country'].value_counts(sort=True,ascending=False,normalize=True)
```

Out[121]:

United-States	0.895857
Mexico	0.019748
Unknown	0.017905
Philippines	0.006081
Germany	0.004207
Canada	0.003716
Puerto-Rico	0.003501
El-Salvador	0.003255
India	0.003071
Cuba	0.002918
England	0.002764
Jamaica	0.002488
South	0.002457
China	0.002303
Italy	0.002242
Dominican-Republic	0.002150
Vietnam	0.002058
Guatemala	0.001966
Japan	0.001904
Poland	0.001843
Columbia	0.001812
Taiwan	0.001566
Haiti	0.001351
Iran	0.001321
Portugal	0.001136
Nicaragua	0.001044
Peru	0.000952
France	0.000891
Greece	0.000891
Ecuador	0.000860
Ireland	0.000737
Hong	0.000614
Trinidad&Tobago	0.000584
Cambodia	0.000584
Thailand	0.000553
Laos	0.000553
Yugoslavia	0.000491
Outlying-US(Guam-USVI-etc)	0.000430
Hungary	0.000399
Honduras	0.000399
Scotland	0.000369
Holand-Netherlands	0.000031

Name: Native-country, dtype: float64

Now let us split categories based on whether or not it is US


In [122]:

```
def func(x):  
    if x!=' United-States':  
        return ' Non-US'  
    else:  
        return x  
df['Native-country']=df['Native-country'].apply(func)
```

In [123]:

```
x1='Native-country'  
y='salary'  
dfp=(df  
.groupby(x1)[y]  
.value_counts(normalize=True)  
.mul(100)  
.rename('percent')  
.reset_index())  
plt.figure(figsize=[15,7])  
g=sns.catplot(x=x1,y='percent',hue=y,data=dfp,kind='bar',height=5,aspect=2)  
for container in g.ax.containers:  
    g.ax.bar_label(container, fmt='%.2f', padding=2)  
plt.title('Native country salary plot')  
plt.savefig('Nativecplot.png')
```

<Figure size 1080x504 with 0 Axes>



Feature Extraction

In [124]:



```
X=df.drop(columns=['fnlwtg', 'salary', 'education', 'work class'])# We are choosing categories
# could clean upto binary fe
#dropping other input columns

X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  category
1   education-num         32561 non-null  category
2   marital status       32561 non-null  category
3   Occupation            32561 non-null  category
4   relationship          32561 non-null  category
5   race                  32561 non-null  category
6   sex                   32561 non-null  category
7   capital gain          32561 non-null  category
8   capital loss          32561 non-null  category
9   Hours-per-week        32561 non-null  category
10  Native-country        32561 non-null  object
dtypes: category(10), object(1)
memory usage: 573.7+ KB
```

In [125]:



```
y=df['salary']
```

Label Encoding for all input categories

In [126]:



```
from sklearn.preprocessing import LabelEncoder
enc=LabelEncoder()
```

In [127]:

```
X1=pd.DataFrame()
for i in list(X.columns):
    X1[i]=enc.fit_transform(X[i])
X1.head()
```

Out[127]:

	age	education- num	marital status	Occupation	relationship	race	sex	capital gain	capital loss	Hours- per- week	Nativ count
0	0	1	1	1	0	1	1	0	0	0	
1	1	1	0	0	1	1	1	0	0	0	
2	0	0	1	1	0	1	1	0	0	0	
3	1	0	0	1	1	0	1	0	0	0	
4	0	1	0	0	1	0	0	0	0	0	

Train test Split

In [128]:

```
from sklearn.model_selection import train_test_split
```

In [129]:

```
X_train, X_test, y_train, y_test = train_test_split(X1, y, test_size=0.1, random_state=42)
```

Train

In [130]:

```
from sklearn.naive_bayes import BernoulliNB
```

In [131]:

```
Bnb=BernoulliNB()
```

In [132]:

```
Bnb.fit(X_train,y_train)
ypred=Bnb.predict(X_test)
```

Performance

In [133]:



```
from sklearn.metrics import accuracy_score, confusion_matrix
```

In [134]:



```
accuracy_score(y_test, ypred)
```

Out[134]:

```
0.7985876573533927
```

In [135]:



```
confusion_matrix(y_test, ypred)
```

Out[135]:

```
array([[1995,  461],
       [ 195,  606]], dtype=int64)
```